



John Dugger / xv6-pa4



Compare

Merge



master

change destination

PA4

change source



Diff Commits Merged pull requests

Files changed (17)

+1	-0	M	.gitignore
+2	-1	M	Makefile
+21	-0	A	Vagrantfile
+1	-0	M	defs.h
+4	-0	M	fcntl.h
+0	-1	M	file.c
+88	-5	M	fs.c
+5	-0	M	fs.h
+154	-0	A	fsTest.c
+5	-1	M	ls.c
+9	-0	M	stat.h
+2	-0	M	syscall.c
+1	-0	M	syscall.h
+29	-5	M	sysfile.c
+1	-0	M	user.h
+1	-0	M	usys.S
+8	-0	A	vagrantprov.sh

.gitignore MODIFIED

Side-by-side diff

View file



```
14 14 kernelmemfs
15 15 mkfs
16 16 .gdbinit
17 +.vagrant
```

Makefile MODIFIED

Side-by-side diff

View file



```
173 173     _stressfs\
174 174     _usertests\
175 175     _wc\
176 +   _fsTest\
176 177     _zombie\
177 178
178 179 fs.img: mkfs README $(UPROGS)
```



```
243 244 EXTRA=\
244 245     mkfs.c ulib.c user.h cat.c echo.c forktest.c grep.c kill.c\
245 246     ln.c ls.c mkdir.c rm.c stressfs.c usertests.c wc.c zombie.c\
246 -   printf.c umalloc.c\
247 +   printf.c umalloc.c fsTest.c\
247 248     README dot-bochsrc *.pl toc.* runoff runoff1 runoff.list\
248 249     .gdbinit.tmpl gdbutil\
249 250
```





Vagrantfile **ADDED**

Side-by-side diff

View file



```
1  +# -*- mode: ruby -*-
2  +# vi: set ft=ruby :
3  +
4  +# Vagrantfile API/syntax version. Don't touch unless you know what you're doing!
5  +VAGRANTFILE_API_VERSION = "2"
6  +
7  +Vagrant.configure(VAGRANTFILE_API_VERSION) do |config|
8  +  config.vm.box = "ubuntu/trusty64"
9  +
10 +  # Share an additional folder to the guest VM. The first argument is
11 +  # the path on the host to the actual folder. The second argument is
12 +  # the path on the guest to mount the folder. And the optional third
13 +  # argument is a set of non-required options.
14 +  # config.vm.synced_folder "../data", "/vagrant_data"
15 +
16 +  config.vm.provider "virtualbox" do |vb|
17 +    vb.customize ["modifyvm", :id, "--memory", "1024"]
18 +  end
19 +
20 +  config.vm.provision :shell, path: "vagrantprov.sh"
21 +end
```

defs.h **MODIFIED**

Side-by-side diff

View file



```
33 33  int          fileread(struct file*, char*, int n);
34 34  int          filestat(struct file*, struct stat*);
35 35  int          filewrite(struct file*, char*, int n);
36 +int          lseek(int fd, int offset);
36 37
37 38  // fs.c
38 39  void          readsb(int dev, struct superblock *sb);
```

fcntl.h **MODIFIED**

Side-by-side diff

View file



```
2 2  #define O_WRONLY 0x001
3 3  #define O_RDWR 0x002
4 4  #define O_CREATE 0x200
5  +// add an O_EXTENT flag to the open() system call
6  +// that will create an extent based file
7  +// the number 0x201 is arbitrary
8  +#define O_EXTENT 0x201
```

file.c **MODIFIED**

Side-by-side diff

View file



```
154 154  }
155 155  panic("filewrite");
156 156  }
157 -
```

fs.c **MODIFIED**

Side-by-side diff

View file



```
370 370
371 371  // Return the disk block address of the nth block in inode ip.
372 372  // If there is no such block, bmap allocates one.
373 +
373 374  static uint
374 375  bmap(struct inode *ip, uint bn)
```



```
375 376 {
376 377     uint addr, *a;
377 378     struct buf *bp;
378
379 + // PA4 changes: Needs to be modified to allocate memory in an extent format
380 +if (ip->type == T_EXTENT) {
381 +     uint blo = 0;
382 +     int i;
383 +     for (i = 0; i < NDIRECT + 1; i++) {
384 +         // (ptr,length)
385 +         int blockPtr = (ip->addrs[i] >> 8); //ptr for first block
386 +         int extentLen = (ip->addrs[i] & 0xff); //length of extent
387 +         if (extentLen == 0) { // Block not currently allocated
388 +             if (i == NDIRECT) {
389 +                 int addr = balloc(ip->dev);
390 +                 blockPtr = (ip->addrs[NDIRECT] >> 8);
391 +                 extentLen = (ip->addrs[NDIRECT] & 0xff);
392 +                 // Block adjacent to last pair i.e its the neighbouring block, we need to increase the last extent
393 +                 if (blockPtr + extentLen == addr) {
394 +                     ip->addrs[NDIRECT] = (blockPtr << 8) | (extentLen + 1);
395 +                 } else {
396 +                     panic("bmap: file exceeds maximum extents");
397 +                 }
398 +             }
399 +             // if not adjucent to the last pair, need to allocate a new pair
400 +             else {
401 +                 int addr = balloc(ip->dev);
402 +                 if (bn == 0 && extentLen == 0) {
403 +                     // allocating first pair and returning its address
404 +                     ip->addrs[0] = (addr << 8) | 1;
405 +                     return addr;
406 +                 } else {
407 +                     // adding another pair if this is not the first pair
408 +                     // resetting values of first block ptr and extent length/size
409 +                     blockPtr = (ip->addrs[bn - 1] >> 8);
410 +                     extentLen = (ip->addrs[bn - 1] & 0xff);
411 +                     // Now checking again if the block is adjucent to last block and coalesce it into pair
412 +                     // if not adjucent , creating new pair in else->
413 +                     if (blockPtr + extentLen == addr) {
414 +                         ip->addrs[i - 1] = (blockPtr << 8) | (extentLen + 1);
415 +                         return addr;
416 +                     } else {
417 +                         ip->addrs[i] = (addr << 8) | 1;
418 +                         return addr;
419 +                     }
420 +                 }
421 +             }
422 +         }
423 +         // Returning block from this pair, if not found check next pair
424 +         if (bn < blo + extentLen) {
425 +             uint offset = bn - blo;
426 +             return blockPtr + offset;
427 +         } else {
428 +             blo = blockPtr + extentLen;
429 +         }
430 +     }
431 + }
432 + //For Direct block
379 433 if(bn < NDIRECT){
380 434     if((addr = ip->addrs[bn]) == 0)
381 435         ip->addrs[bn] = addr = balloc(ip->dev);
382 436     return addr;
383 437 }
384 438 bn -= NDIRECT;
385
439 +// for indirect block
386 440 if(bn < NINDIRECT){
387 441     // Load indirect block, allocating if necessary.
388 442     if((addr = ip->addrs[NDIRECT]) == 0)
396 450         brelse(bp);
397 451     return addr;
398 452 }
```





```
399 -
400 453     panic("bmap: out of range");
401 454 }
402 455
...
411 464     int i, j;
412 465     struct buf *bp;
413 466     uint *a;
414 -
467 + //getting length of extent
468 + if (ip->type == T_EXTENT) {
469 +     int i;
470 +     for (i = 0; i < NDIRECT + 1; i++) {
471 +         // (ptr, length)pair
472 +         int blockPtr = (ip->addrs[i] >> 8);
473 +         int extentLen = (ip->addrs[i] & 0xff);
474 +         if (extentLen != 0) {
475 +             int j;
476 +             for(j = 0; j < extentLen; j++) { bfree(ip->dev, blockPtr+j); }
477 +         }
478 +     }
479 + }
480 + else {
415 481     for(i = 0; i < NDIRECT; i++){
416 482         if(ip->addrs[i]){
417 483             bfree(ip->dev, ip->addrs[i]);
...
430 496         bfree(ip->dev, ip->addrs[NDIRECT]);
431 497         ip->addrs[NDIRECT] = 0;
432 498     }
433 -
499 + }
434 500     ip->size = 0;
435 501     iupdate(ip);
436 502 }
...
440 506 void
441 507 stati(struct inode *ip, struct stat *st)
442 508 {
509 + // fstat() uses stati to dump information. Need to also
510 + // modify struct stat in stat.h to hold the information
511 + // modify fstat() system call such that it will dump
512 + // information about each extent of an extent based
513 + // file in addition to file size
443 514     st->dev = ip->dev;
444 515     st->ino = ip->inum;
445 516     st->type = ip->type;
446 517     st->nlink = ip->nlink;
447 518     st->size = ip->size;
519 + // TA directed returning the address even though default
520 + // fstat() doesn't return the address. Added
521 + int i;
522 + for (i = 0; i < NDIRECT + 1; i++) {
523 +     st->bloadrs[i] = ip->addrs[i];
524 + }
525 +
526 +
448 527 }
449 528
450 529 //PAGEBREAK!
...
464 543
465 544     if(off > ip->size || off + n < off)
466 545         return -1;
546 + // if offset + > maxfilesize*block and type is not equal to T_EXTENT:
547 + //limits the size of each extent to 2^8 blocks and the disk addresses to 2^24.
548 + //if(off + n > MAXFILE*BSIZE && ip -> type != T_EXTENT)
549 + // n = MAXFILE*BSIZE - off;
467 550     if(off + n > ip->size)
468 551         n = ip->size - off;
469 552
```

fs.h

MODIFIED

Side-by-side diff

View file

...

```
24 24 #define NDIRECT 12
25 25 #define NINDIRECT (BSIZE / sizeof(uint))
26 26 #define MAXFILE (NDIRECT + NINDIRECT)
27 27 + #define MAXEXTENT (6 * 256)
27 28
28 29 // On-disk inode structure
29 30 struct dinode {
32 33     short minor;           // Minor device number (T_DEV only)
33 34     short nlink;          // Number of links to inode in file system
34 35     uint size;            // Size of file (bytes)
36 36 + // keep the inode structure exactly as it
37 37 + // is except for those data block points.
38 38 + // Use 3 of the 4 bytes for pointer and
39 39 + // the remaining byte for length.
35 40     uint addrs[NDIRECT+1]; // Data block addresses
36 41 };
37 42
```

fsTest.c

ADDED

Side-by-side diff

View file

...

```
1 + #include "types.h"
2 + #include "stat.h"
3 + #include "user.h"
4 + #include "fs.h"
5 + #include "fcntl.h"
6 + #include "syscall.h"
7 + #include "traps.h"
8 +
9 + int stdout = 1;
10 +
11 + int
12 + main (int argc, char* argv[]) {
13 +     int fd = open("mytestFile.txt", O_CREATE | O_RDWR | O_EXTENT);
14 +     int fd2 = open("mytestFile2.txt", O_CREATE | O_RDWR | O_EXTENT);
15 +     struct stat st, st2; //for fstat
16 +     if (fd <= 1) {
17 +         printf(stdout, "Issue creating file1\n");
18 +     }
19 +
20 +     if (fd2 <= 1) {
21 +         printf(stdout, "Issue creating file2\n");
22 +     }
23 +     printf(stdout, "***** Starting test for file1 *****\n");
24 +     int i;
25 +     for (i = 0; i < 65; i++) {
26 +         if (write(fd, "abcdef\n", 8) != 8) {
27 +             printf(stdout, "Error: Writing abcdef to new file1 line %d failed\n", i);
28 +             exit();
29 +         }
30 +         else { printf(stdout, "Writing abcdef to line%d of file1\n", i); }
31 +     }
32 +
33 +     fstat(fd, &st);
34 +     printf(stdout, "File size i= %d\n", st.size);
35 +     printf(stdout, "File type = %d\n", st.type);
36 +     printf(stdout, "File system device = %d\n", st.dev);
37 +     printf(stdout, "inode num = %d\n", st.ino);
38 +     printf(stdout, "Number of links = %d\n", st.nlink);
39 +     printf(stdout, "Block Addresses:\n");
40 +     if (st.type == T_EXTENT) {
41 +         for (i = 0; i < NDIRECT + 1; i++) {
42 +             //first block pointer
43 +             int bloPtr = (st.bloaddrs[i] >> 8);
```



```
44 +     int len = (st.bloaddrs[i] & 0xff); // length of extent
45 +     printf(stdout, "   Pointer %d's address = %d with size = %d\n", i, bloPtr, len);
46 + }
47 + } else {
48 +     for (i = 0; i < NDIRECT + 1; i++) {
49 +         printf(stdout, "   Pointer %d's address = %d\n", i, st.bloaddrs[i]);
50 +     }
51 + }
52 + //printf(stdout, "Disk address of block = %d\n", st.bloaddrs);
53 + printf(stdout, "Number of blocks = %d\n", st.length);
54 + if (st.size != (512 + 8) ) {
55 +     printf(stdout, "File size is different from expected, size = %d\n", st.size);
56 +     printf(stdout, "File type = %d\n", st.type);
57 + }
58 +
59 + printf(stdout, "***** Starting test for file2 *****\n");
60 + int j;
61 + for (j = 0; j < 65; j++) {
62 +     if(write(fd2, "hello\n", 7) != 7) {
63 +         printf(stdout, "Error: Writing hello to file2 line%d failed\n", i);
64 +         exit();
65 +     }
66 +     else{printf(stdout, "Writing hello to line%d of file2\n", j);}
67 + }
68 +
69 + fstat(fd2, &st2);
70 + printf(stdout, "File size i= %d\n", st2.size);
71 + printf(stdout, "File type = %d\n", st2.type);
72 + printf(stdout, "File system device = %d\n", st2.dev);
73 + printf(stdout, "inode num = %d\n", st2.ino);
74 + printf(stdout, "Number of links = %d\n", st2.nlink);
75 + printf(stdout, "Block Addresses:\n");
76 + if (st2.type == T_EXTENT) {
77 +     for (j = 0; j < NDIRECT + 1; j++) {
78 +         //first block pointer
79 +         int bloPtr2 = (st2.bloaddrs[j] >> 8);
80 +         int len2 = (st2.bloaddrs[j] & 0xff); // length of extent
81 +         printf(stdout, " File2: Pointer %d's address = %d with size = %d\n", j, bloPtr2, len2);
82 +     }
83 + } else {
84 +     for (j = 0; j < NDIRECT + 1; j++) {
85 +         printf(stdout, " File2: Pointer %d's address = %d\n", j, st2.bloaddrs[j]);
86 +     }
87 + }
88 + //printf(stdout, "Disk address of block = %d\n", st.bloaddrs);
89 + printf(stdout, "Number of blocks = %d\n", st2.length);
90 + if (st2.size != (448 + 7) ) {
91 +     printf(stdout, "File2 size is different from expected, size = %d\n", st2.size);
92 +     printf(stdout, "File2 type = %d\n", st2.type);
93 + }
94 + //lseek testing begins here
95 + int fd3;
96 + if(fd3 = open("mytestFile3.txt", O_CREATE | O_RDWR | O_EXTENT) <= 1){
97 +     printf(stdout, "Error: Issue creating file1\n");
98 + }
99 + if(write(fd3, "abcdefghijklmnopqrstuvwxyz0123456789ABCDEFGHIJKLMNOPQRSTUVWXYZ\n", 64) != 64){
100 +     printf(stdout, "Error: Writing to file3 failed\n");
101 +     exit();
102 + }
103 + printf(stdout, "Testing lseek(). Current file contents:\n");
104 +
105 + // output contents of the file
106 + int n;
107 + char buf[512];
108 + while((n = read(fd3, buf, sizeof(buf))) > 0) {
109 +     if (write(1, buf, n) != n) {
110 +         printf(1, "cat: write error\n");
111 +         exit();
112 +     }
113 + }
114 + if(n < 0)
115 +     printf(1, "cat: read error\n");
116 +
117 + // the offset should be 64 before this
```



```
118 + if (lseek(fd3, 26) == -1){
119 +     printf(stdout, "lseek error\n");
120 + }
121 + if(write(fd3, "ABCDEFGHJKLMNOPQRSTUVWXYZ0123456789", 36) != 36) {
122 +     printf(stdout, "Error: Writing to file3 failed\n");
123 +     exit();
124 + }
125 + // should see a gap of 10 zeroes before the last hello\n
126 + printf(stdout, "Capitals and digits swapped. Current file contents:\n");
127 +
128 + // output contents of the file
129 + while((n = read(fd, buf, sizeof(buf))) > 0) {
130 +     if (write(1, buf, n) != n) {
131 +         printf(1, "cat: write error\n");
132 +         exit();
133 +     }
134 + }
135 + if(n < 0)
136 +     printf(1, "cat: read error\n");
137 +
138 + if(lseek(fd, -1) == 0){
139 +     printf(stdout, "Error: lseek() allowed negative offset\n");
140 + } else{
141 +     printf(stdout, "lseek() correctly didn't allow a negative offset\n");
142 + }
143 + if(lseek(fd, 0x7FFFFFFF) == 0){
144 +     printf(stdout, "Error: lseek() allowed offset outside the end of the file\n");
145 + } else{
146 +     printf(stdout, "lseek() correctly didn't allow offset outside the end of the file\n");
147 + }
148 +
149 + printf(stdout, "*****Testing Extent File system completed!*****\n");
150 + close(fd);
151 + close(fd2);
152 + close(fd3);
153 + exit();
154 +}
```



ls.c MODIFIED

Side-by-side diff

View file



```
42 42     }
43 43
44 44     switch(st.type){
45 45 - case T_FILE:
46 46 + case T_FILE: // or T_EXTENT
47 47 +     printf(1, "%s %d %d %d\n", fmtname(path), st.type, st.ino, st.size);
48 48 +     break;
49 49 + case T_EXTENT: // handled exactly like T_FILE
50 50     printf(1, "%s %d %d %d\n", fmtname(path), st.type, st.ino, st.size);
51 51     break;
52 52
```



stat.h

MODIFIED

Side-by-side diff

View file



```
1 1 #define T_DIR 1 // Directory
2 2 #define T_FILE 2 // File
3 3 #define T_DEV 3 // Device
4 4 +// create a new type of file to support
5 5 +// inodes with pointers and length
6 6 +#define T_EXTENT 4 // Extent based file
7 7
8 8 struct stat {
9 9     short type; // Type of file
10 10
11 11     uint ino; // Inode number
12 12     short nlink; // Number of links to file
```





```
10 13      uint size;    // Size of file in bytes
14 + // adding an address after discussion with the TA
15 + // address wasn't included for the default fstat
16 + uint bloaddrs[13]; // Disk address of block
17 + // length will be 1 by default for T_FILE /T_DIR / T_DEV types
18 + // length will be the last byte of the address for T_EXTENT
19 + uint length; // Number of blocks
11 20  };
```

syscall.c **MODIFIED**

Side-by-side diff

View file



```
103 103      extern int sys_wait(void);
104 104      extern int sys_write(void);
105 105      extern int sys_uptime(void);
106 +extern int sys_lseek(void);
106 107
107 108      static int (*syscalls[])(void) = {
108 109          [SYS_fork]    sys_fork,
```

```
126 127          [SYS_link]    sys_link,
127 128          [SYS_mkdir]    sys_mkdir,
128 129          [SYS_close]    sys_close,
130 +[SYS_lseek]    sys_lseek,
129 131      };
130 132
131 133      void
```

syscall.h **MODIFIED**

Side-by-side diff

View file



```
20 20      #define SYS_link    19
21 21      #define SYS_mkdir    20
22 22      #define SYS_close    21
23 +#define SYS_lseek    22
```

sysfile.c **MODIFIED**

Side-by-side diff

View file



```
248 248      if((dp = nameiparent(path, name)) == 0)
249 249          return 0;
250 250      ilock(dp);
251 -
251 +// checking if file is existing
252 252      if((ip = dirlookup(dp, name, &off)) != 0){
253 253          iunlockput(dp);
254 254          ilock(ip);
255 - if(type == T_FILE && ip->type == T_FILE)
255 + if((type == T_FILE && ip->type == T_FILE) || (type == T_EXTENT && ip->type == T_EXTENT))
256 256          return ip;
257 257          iunlockput(ip);
258 258          return 0;
259 259      }
260 260
261 +// creating file coz it doesnt exist
261 262      if((ip = ialloc(dp->dev, type)) == 0)
262 263          panic("create: ialloc");
263 264
```

```
297 298      begin_op();
298 299
299 300      if(omode & O_CREATE){
300 - ip = create(path, T_FILE, 0, 0);
301 - if(ip == 0){
302 - end_op();
```





```
301 +   if (omode & O_EXTENT) {
302 +       if((ip = create(path, T_EXTENT, 0, 0)) == 0)
303 +           return -1;
304 +   }
305 +   else {
306 +       if((ip = create(path, T_FILE, 0, 0)) == 0)
303 307           return -1;
304 308   }
305 309   } else {
    ...
443 447   fd[1] = fd1;
444 448   return 0;
445 449   }
450 +
451 +int
452 +sys_lseek(void)
453 +{
454 +   int fd;
455 +   int offset;
456 +   struct file *f;
457 +
458 +   // checks to verify that there are valid arguments
459 +   // stores the arguments in fd, f and offset
460 +   // makes sure the offset is within the file
461 +   if(argfd(0, &fd, &f) < 0 || argint(1, &offset) < 0 || offset < 0 || offset > f->ip->size)
462 +   {
463 +       return -1;
464 +   }
465 +   // f now stores the file referenced by fd
466 +   // changes the offset for the file to the offset value passed in
467 +   f->off = offset;
468 +   return f->off;
469 +}
```



user.h

MODIFIED

Side-by-side diff

View file



```
23 23   char* sbrk(int);
24 24   int sleep(int);
25 25   int uptime(void);
26   +int lseek(int, int);
26 27
27 28   // ulib.c
28 29   int stat(char*, struct stat*);
    ...
```



usys.S

MODIFIED

Side-by-side diff

View file



```
29 29   SYSCALL(sbrk)
30 30   SYSCALL(sleep)
31 31   SYSCALL(uptime)
32   +SYSCALL(lseek)
```



vagrantprov.sh

ADDED

Side-by-side diff

View file



```
1   +#!/usr/bin/env bash
2   +
3   +apt-get update
4   +apt-get install -y qemu-system-x86
5   +apt-get install -y gdb
6   +apt-get install -y tmux
7   +
8   +echo "Set auto-load safe-path /" > /home/vagrant/.gdbinit
```



